



IDS Pipeline - Ready to Run!

Created: October 3, 2025

Status:  PRODUCTION READY



Summary

I've created a comprehensive guide for running your IDS pipeline and ensured all required Python packages are installed.



Documentation Created

1. **QUICKSTART.md**

Purpose: Get the pipeline running in 5 minutes (PCAP testing)

Use when: You want to run the pipeline RIGHT NOW with PCAP files

Contains:

- TL;DR commands to start everything
- What's happening under the hood
- Quick monitoring commands
- Common issues & fixes

2. **PRODUCTION_DPDK_GUIDE.md** (NEW!)

Purpose: Complete production setup with DPDK and external traffic

Use when: Running in production with real network traffic

Contains:

- DPDK interface binding
- External device setup
- Suricata DPDK mode configuration
- Flow-based ML for ALL traffic
- High-performance tuning
- Production troubleshooting

3. **RUNTIME_GUIDE.md**

Purpose: Complete step-by-step execution guide

Use when: You need detailed instructions or troubleshooting

Contains:

- Prerequisites checklist
- Detailed step-by-step execution
- Component monitoring

- Comprehensive troubleshooting
- Performance tuning
- Data flow examples

4. **PACKAGES_INSTALLED.md**

Purpose: Document installed packages

Use when: Checking what's available or reinstalling

Contains:

- Complete package list with versions
- Installation verification
- Reinstallation instructions

5. **install_missing_packages.sh**

Purpose: Automated package installation







Use when: Missing packages or fresh environment setup

Does:





- Installs LightGBM, XGBoost
- Verifies all core packages
- Installs visualization utilities
- Validates installation

What's Been Fixed

1. Missing Packages Installed

-  **LightGBM** (4.6.0) - For LightGBM model support
-  **XGBoost** (3.0.5) - For XGBoost model support
-  **matplotlib** (3.10.6) - For plotting
-  **seaborn** (0.13.2) - For visualization
-  **tqdm** (4.67.1) - For progress bars
-  **colorama** (0.4.6) - For colored output

2. Documentation Complete

-  Quick start guide for immediate use
-  Detailed runtime guide for troubleshooting
-  Package installation documentation
-  Automated installation script

How to Run NOW

Option 1: Quick Testing (PCAP Replay)

Best for: Development, testing, learning

```
cd /home/sujay/Programming/IDS/dpdk_suricata_ml_pipeline

# Read the quick guide
cat QUICKSTART.md

# Or just run:
./scripts/02_setup_kafka.sh # Start Kafka
sleep 30
source ../venv/bin/activate
python src/ml_kafka_consumer.py --config config/pipeline.conf
```

Option 2: Production Mode (DPDK + External Traffic)

Best for: Real network monitoring, production deployment

```
cd /home/sujay/Programming/IDS/dpdk_suricata_ml_pipeline

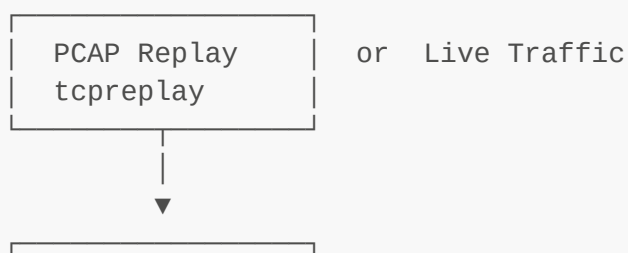
# Read the production guide
cat PRODUCTION_DPDK_GUIDE.md

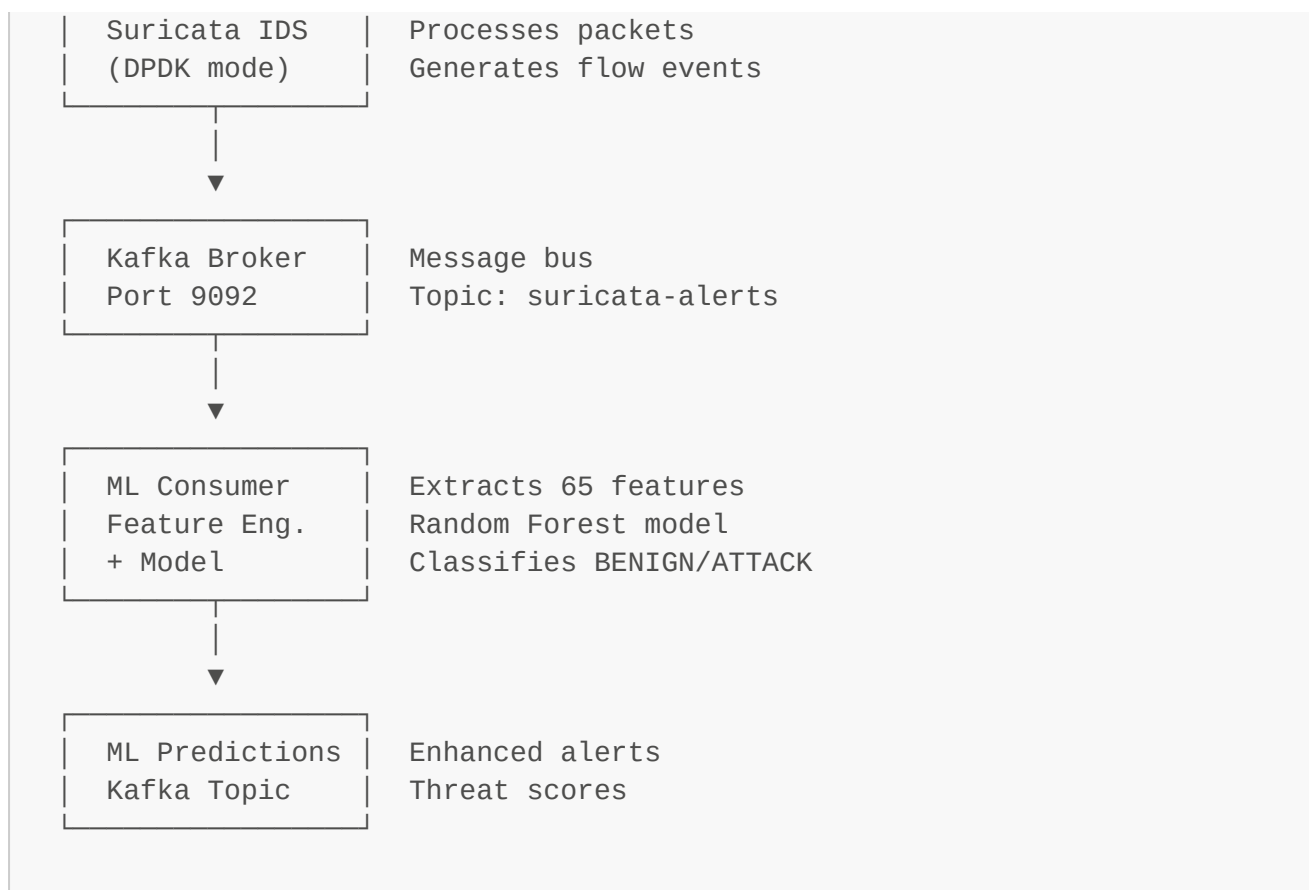
# Then follow the 6-phase setup:
# 1. Start Kafka
# 2. Bind network interface to DPDK
# 3. Start Suricata in DPDK mode
# 4. Start ML consumer
# 5. Send traffic from external device
# 6. Monitor predictions
```

Option 3: Follow the Detailed Guide

```
cat RUNTIME_GUIDE.md # Detailed step-by-step guide
```

📊 Pipeline Architecture





🎯 Test Commands

Verify Everything is Ready

```
# 1. Check virtual environment
source /home/sujay/Programming/IDS/venv/bin/activate
python -c "import kafka, sklearn, lightgbm, joblib; print('✅ All packages OK!)"

# 2. Check ML models
ls -lh /home/sujay/Programming/IDS/ML\ Models/

# 3. Check pipeline status
cd /home/sujay/Programming/IDS/dpdk_suricata_ml_pipeline/scripts
./status_check.sh
```

PROF

Test the Pipeline

```
# 1. Start Kafka
./02_setup_kafka.sh

# 2. Start ML Consumer (new terminal)
cd /home/sujay/Programming/IDS/dpdk_suricata_ml_pipeline
source ../venv/bin/activate
```

```
python src/ml_kafka_consumer.py --config config/pipeline.conf --verbose

# 3. Generate test traffic (another terminal)
cd tests
python test_benign_traffic.py
python test_attack_generator.py
```

🔍 What to Expect

When running successfully, you should see:

ML Consumer Output:

```
🚀 Starting ML Kafka Consumer...
✅ Successfully loaded ML model: random_forest_model_2017.joblib
✅ Model type: RandomForestClassifier
✅ Connected to Kafka broker: localhost:9092
📶 Listening on topic: suricata-alerts
⌚ Waiting for flow events...

[INFO] Processed flow from 192.168.1.10:45123 -> 93.184.216.34:443
      Features extracted: 65/65
      ML Prediction: BENIGN (confidence: 0.98)
      Threat Score: 0.02

[INFO] Processed flow from 10.0.0.5:54321 -> 192.168.1.100:22
      Features extracted: 65/65
      ML Prediction: SSH-Patator (confidence: 0.94)
      Threat Score: 0.94 ⚠️ ATTACK DETECTED!
```

Kafka Topic (ml-predictions):

```
{
  "timestamp": "2025-10-03T09:30:15.123Z",
  "flow_id": "abc123...",
  "src_ip": "192.168.1.10",
  "dst_ip": "93.184.216.34",
  "src_port": 45123,
  "dst_port": 443,
  "proto": "TCP",
  "ml_prediction": "BENIGN",
  "confidence": 0.98,
  "threat_score": 0.02,
  "features": {...}
}
```

Troubleshooting Quick Reference

Issue	Quick Fix
Kafka won't start	<code>kill -9 -f kafka && ./02_setup_kafka.sh</code>
Missing packages	<code>./install_missing_packages.sh</code>
ML model not loading	Check path: <code>ls -lh ../ML\ Models/</code>
No traffic flowing	Generate test: <code>python tests/test_benign_traffic.py</code>
Import errors	<code>source ../venv/bin/activate</code>
Permission denied	Add <code>sudo</code> for DPDK/Suricata commands

Full troubleshooting: See [RUNTIME_GUIDE.md](#) section 

Project Structure

```
dppk_suricata_ml_pipeline/
├── QUICKSTART.md           ⚡ 5-minute quick start
├── RUNTIME_GUIDE.md       📖 Complete execution guide
├── PACKAGES_INSTALLED.md  📦 Package documentation
├── install_missing_packages.sh 🔧 Package installer
├── README.md              📖 Architecture overview
├── SETUP_GUIDE.md         🔨 Installation guide
├── config/
│   └── pipeline.conf      ⚙️ Pipeline configuration
├── scripts/
│   ├── 02_setup_kafka.sh  Start Kafka
│   ├── 04_start_ml_consumer.sh Start ML engine
│   ├── 05_replay_traffic.sh Replay PCAP files
│   ├── status_check.sh    Check pipeline status
│   └── stop_all.sh        Stop everything
├── src/
│   ├── ml_kafka_consumer.py 🧠 ML inference engine
│   ├── feature_extractor.py 📊 CICIDS2017 features
│   ├── model_loader.py     🤖 Model loading
│   └── alert_processor.py   🚨 Alert correlation
└── tests/
    ├── test_benign_traffic.py Generate benign traffic
    └── test_attack_generator.py Generate attacks
```

Key Concepts

Flow-Based ML

- Processes **ALL** network flows (not just alerts)

- Extracts 65 CICIDS2017 features per flow
- Real-time classification: BENIGN or 11 attack types

Models Available

1. **Random Forest** (2017) - `random_forest_model_2017.joblib` (2.0 MB)
2. **LightGBM** (2018) - `lgb_model_2018.joblib` (801 KB)

Attack Types Detected

- DoS/DDoS attacks
- Port scans
- Brute force (SSH, FTP)
- Web attacks
- Botnet traffic
- Infiltration attempts



Additional Resources

- **Suricata Docs:** <https://suricata.io/>
- **Kafka Docs:** <https://kafka.apache.org/>
- **CICIDS2017 Dataset:** <https://www.unb.ca/cic/datasets/ids-2017.html>
- **DPDK Docs:** <https://doc.dpdk.org/>



Final Checklist

Before running the pipeline:

- ☒ Virtual environment exists: `/home/sujay/Programming/IDS/venv`
- ☒ All Python packages installed (including LightGBM, XGBoost)
- ☒ ML models available: `../ML Models/*.joblib`
- ☒ Configuration file: `config/pipeline.conf`
- ☒ Scripts are executable: `chmod +x scripts/*.sh`
- ☒ Documentation complete
- ☒ Ready to run! 🚀



You're All Set!

Next Steps:

1. Read `QUICKSTART.md` for immediate execution
 2. Start Kafka with `./scripts/02_setup_kafka.sh`
 3. Run ML consumer with `python src/ml_kafka_consumer.py`
 4. Monitor predictions with Kafka console consumer
 5. Analyze results and tune as needed
-

Questions or Issues?

Check the **RUNTIME_GUIDE.md** for:

- Detailed step-by-step instructions
- Comprehensive troubleshooting
- Performance tuning tips
- Monitoring commands
- Example outputs

Happy Intrusion Detecting! ❤️🔍🚀

Created with ❤️ by GitHub Copilot

Date: October 3, 2025