

DPDK-Suricata-Kafka-ML Complete Pipeline

Architecture Overview

DPDK-Suricata-Kafka-ML Complete Pipeline

Architecture Overview

```
External Traffic (tcpreplay/PCAP)
  ↓ (via Ethernet)
DPDK Packet Capture (bound interface)
  ↓ (zero-copy)
Suricata IDS (DPDK mode)
  ↓ (eve-kafka output: flows + alerts)
Kafka Broker
  ↓ (consume ALL events)
ML Inference Engine (Flow-Based)
  ↓ (CICIDS2017 feature extraction + predictions for EVERY flow)
Enhanced Alert Stream
  ↓
Alert Dashboard / Database
```

Key Features

- 🌟 **Flow-Based ML Inference:** Process **ALL network flows**, not just signature alerts
- 🧠 **CICIDS2017 Feature Extraction:** 65-feature extraction from Suricata flow events
- 🔍 **Dual Detection:** Combines Suricata signature detection + ML anomaly detection
- ⚡ **Real-Time Processing:** Streaming architecture with Kafka message bus
- 📊 **Threat Scoring:** Combined threat scores from multiple detection sources

Components

1. DPDK Packet Ingestion

- Binds network interface to DPDK driver (vfio-pci/uio_pci_generic)
- Zero-copy packet capture
- High-performance packet processing

2. Suricata IDS (Enhanced)

- Runs in DPDK mode
- **NEW:** Logs ALL network flows (not just alerts)
- Processes packets with signature-based detection
- Outputs flows + alerts to Kafka via eve-kafka plugin

3. Kafka Streaming

- Acts as message broker
- Buffers flows and alerts between Suricata and ML
- Enables scalable processing
- Topics: **suricata-alerts** (input), **ml-predictions** (output)

4. ML Inference Engine (Enhanced)

- **NEW:** Consumes ALL flow events from Kafka (not just alerts)
- Extracts 65 CICIDS2017 features from every flow
- Performs real-time ML inference on all traffic
- Correlates ML predictions with Suricata alerts
- Generates enhanced alerts with combined threat scores
- Supports Random Forest and LightGBM models

Directory Structure

```
dpdk_suricata_ml_pipeline/
├── README.md                # This file
├── SETUP_GUIDE.md           # Detailed setup instructions
├── config/                  # Configuration files
│   ├── suricata-dpdk.yaml   # Suricata config with DPDK & Kafka
│   ├── kafka-config.properties # Kafka settings
│   └── pipeline.conf        # Pipeline configuration
├── scripts/                 # Management scripts
│   ├── 01_bind_interface.sh # Bind NIC to DPDK
│   ├── 02_setup_kafka.sh    # Install/configure Kafka
│   ├── 03_start_suricata.sh # Start Suricata in DPDK mode (with flow
logging)
│   ├── 04_start_ml_consumer.sh # Start ML inference
│   ├── 05_replay_traffic.sh    # Replay PCAP files
│   ├── stop_all.sh             # Stop all services
│   ├── status_check.sh        # Check pipeline status
│   └── unbind_interface.sh     # Restore network interface
├── src/                     # Python source code
│   ├── ml_kafka_consumer.py   # ML inference consumer (flow-based)
│   ├── feature_extractor.py   # CICIDS2017 65-feature extraction
│   ├── model_loader.py        # ML model loading (RF, LightGBM)
│   └── alert_processor.py     # Alert correlation & threat scoring
├── pcap_samples/            # Sample PCAP files for testing
│   └── README.md
└── logs/                    # Pipeline logs
    ├── dpdk/
    ├── suricata/
    ├── kafka/
    └── ml/
```

PROF

ML Inference Features

Feature Extraction (CICIDS2017 Compatible)

The pipeline extracts **65 network flow features** from Suricata events:

Flow Statistics: Duration, packet counts, byte counts

Packet Length Stats: Min, max, mean, std (forward & backward)

Inter-Arrival Time (IAT): Mean, std, min, max (flow, fwd, bwd)

TCP Flags: FIN, SYN, RST, PSH, ACK, URG, ECE counts

Header Lengths: Forward & backward header sizes

Packet Rates: Packets/sec, bytes/sec (overall, fwd, bwd)

Protocol Features: Protocol type, port numbers

Derived Features: Down/up ratio, avg segment sizes, active/idle times

Supported ML Models

- **Random Forest** (scikit-learn): `random_forest_model_2017.joblib`
- **LightGBM**: `lgb_model_2018.joblib`
- Models located in: `../ML Models/`

Alert Processing

- Correlates Suricata signature alerts with ML predictions
- Calculates combined threat scores (0-1 scale)
- Assigns threat levels: BENIGN, LOW, MEDIUM, HIGH, CRITICAL
- Outputs enhanced alerts to Kafka topic

Quick Start

Prerequisites

- DPDK installed (run `install_dpdk_suricata.sh`)
- Suricata installed with DPDK support
- Kafka installed and running
- Python virtual environment activated
- ML models in `../ML Models/` directory

Step 1: Configure Interface

Edit `config/pipeline.conf` and set your network interface:

```
NETWORK_INTERFACE="eth0" # Change to your interface
```

Step 2: Bind Interface to DPDK

```
cd dpdk_suricata_ml_pipeline
sudo ./scripts/01_bind_interface.sh
```

Step 3: Start Kafka

```
./scripts/02_setup_kafka.sh
```

Step 4: Start Suricata (with Flow Logging)

```
## Components

### 1. DPDK Packet Ingestion
- Binds network interface to DPDK driver (vfio-pci/uio_pci_generic)
- Zero-copy packet capture
- High-performance packet processing

### 2. Suricata IDS
- Runs in DPDK mode
- Processes packets with signature-based detection
- Outputs alerts to Kafka via eve-kafka plugin

### 3. Kafka Streaming
- Acts as message broker
- Buffers alerts between Suricata and ML
- Enables scalable processing

### 4. ML Inference Engine
- Consumes alerts from Kafka
- Extracts features from network events
- Performs real-time threat classification
- Outputs enhanced predictions

## Directory Structure

...

dpdk_suricata_ml_pipeline/
├── README.md                # This file
├── SETUP_GUIDE.md          # Detailed setup instructions
├── config/                  # Configuration files
│   ├── suricata-dpdk.yaml   # Suricata config with DPDK & Kafka
│   ├── kafka-config.properties # Kafka settings
│   └── pipeline.conf        # Pipeline configuration
├── scripts/                 # Management scripts
│   ├── 01_bind_interface.sh # Bind NIC to DPDK
│   ├── 02_setup_kafka.sh    # Install/configure Kafka
│   ├── 03_start_suricata.sh # Start Suricata in DPDK mode
│   ├── 04_start_ml_consumer.sh # Start ML inference
│   ├── 05_replay_traffic.sh # Replay PCAP files
│   ├── stop_all.sh          # Stop all services
│   └── status_check.sh      # Check pipeline status
```

```

├── unbind_interface.sh      # Restore network interface
├── src/                    # Python source code
│   ├── ml_kafka_consumer.py # ML inference consumer
│   ├── feature_extractor.py # Extract features from Suricata alerts
│   ├── model_loader.py     # Load ML models
│   └── alert_processor.py  # Process and store predictions
├── pcap_samples/          # Sample PCAP files for testing
│   └── README.md
└── logs/                  # Pipeline logs
    ├── dpdk/
    ├── suricata/
    ├── kafka/
    └── ml/
...

```

Quick Start

Prerequisites

- DPDK installed (run `install_dpdk_suricata.sh`)
- Suricata installed with DPDK support
- Kafka installed and running
- Python virtual environment activated

Step 1: Configure Interface

Edit `config/pipeline.conf` and set your network interface:

```

```bash
NETWORK_INTERFACE="eth0" # Change to your interface
...

```

### ### Step 2: Bind Interface to DPDK

```

```bash
cd dpdk_suricata_ml_pipeline
sudo ./scripts/01_bind_interface.sh
...

```

Step 3: Start Kafka

```

```bash
./scripts/02_setup_kafka.sh
...

```

### ### Step 4: Start Suricata

```

```bash
sudo ./scripts/03_start_suricata.sh
...

```

Step 5: Start ML Consumer

```

```bash
./scripts/04_start_ml_consumer.sh
...

```

### ### Step 6: Replay Traffic

```

```bash
sudo ./scripts/05_replay_traffic.sh pcap_samples/sample.pcap

```

```

...

### Monitor Pipeline
```bash
./scripts/status_check.sh
```

### Stop Everything
```bash
sudo ./scripts/stop_all.sh
```

## Traffic Sources

### Option 1: tcpreplay (PCAP files)
Replay captured traffic from PCAP files:
```bash
sudo tcpreplay -i eth0 capture.pcap
```

### Option 2: External System
Send live traffic from another machine:
- Configure second machine to send to monitored interface
- Use hping3, scapy, or actual application traffic

### Option 3: Traffic Generator
Use DPDK pktgen or similar:
```bash
dpdk-pktgen -l 0-3 -n 4 -- -P -m "[1:2].0" -f traffic.lua
```

## Monitoring

### Check DPDK Status
```bash
dpdk-devbind.py --status
grep Huge /proc/meminfo
```

### Check Suricata
```bash
tail -f logs/suricata/suricata.log
tail -f /var/log/suricata/eve.json
```

### Check Kafka
```bash
kafka-console-consumer --bootstrap-server localhost:9092 \
 --topic suricata-alerts --from-beginning
```

### Check ML Consumer
```bash
```

```

```
tail -f logs/ml/ml_consumer.log  
``
```

Performance Tuning

DPDK

- Allocate more hugepages: 4-8GB recommended
- Use CPU isolation: `isolcpus=` in GRUB
- Bind to isolated CPU cores

Suricata

- Increase worker threads
- Tune ring buffers
- Disable unnecessary features
- Use optimized rules

Kafka

- Increase buffer sizes
- Adjust retention policies
- Configure compression
- Tune batch sizes

ML Inference

- Batch predictions
- Use GPU if available
- Cache model in memory
- Parallel processing

Troubleshooting

See `SETUP_GUIDE.md` for detailed troubleshooting steps.

License

Part of IDS Project - October 2025