# IDS Pipeline Architecture Guide

### Table of Contents

- Overview
- AF\_PACKET Pipeline
- DPDK Pipeline
- Comparison
- Component Details

#### Overview

This IDS (Intrusion Detection System) uses a multi-stage pipeline combining Suricata for signature-based detection and Machine Learning for anomaly detection. The system can operate in two modes:

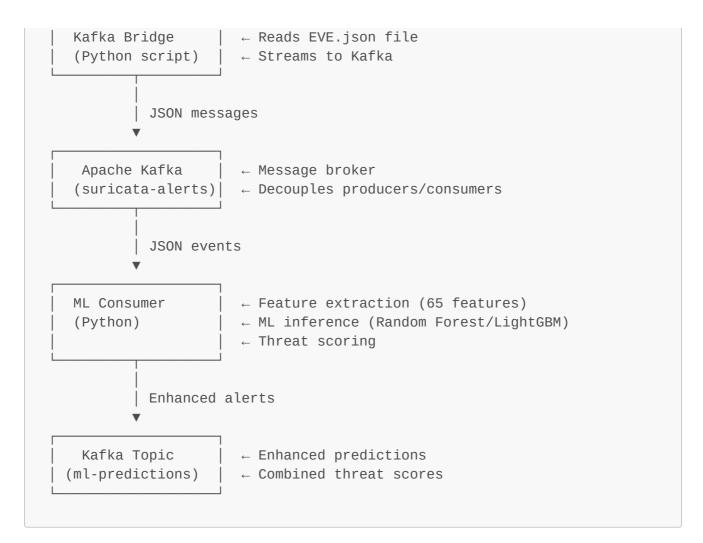
- 1. AF\_PACKET Mode Standard Linux packet capture (compatible with all interfaces)
- 2. DPDK Mode High-performance kernel bypass (requires DPDK-compatible NIC)

Both pipelines follow the same architectural pattern but differ in packet capture mechanisms.

## AF\_PACKET Pipeline

Architecture Flow

```
AF_PACKET MODE PIPELINE
Network Traffic
Network Interface
                    ← Any interface (USB, PCIe, WiFi, etc.)
 (AF_PACKET)
         Raw Packets
  Suricata
                    ← Signature-based detection
                    ← Rule-based alerts
 (AF_PACKET)
         EVE JSON logs (flows + alerts)
```



Now It Works

#### Step 1: Packet Capture (AF\_PACKET)

```
# Interface setup
ip link set enx00e04c36074c up
ip link set enx00e04c36074c promisc on
```

PROF

**AF\_PACKET** is a Linux socket type that captures packets at Layer 2:

- Pros: Works with ANY network interface (USB adapters, WiFi, PCIe NICs)
- Cons: Packets traverse kernel network stack (some overhead)
- **Performance**: Good for moderate traffic (< 1 Gbps)

#### Step 2: Suricata Detection

```
suricata --af-packet=enx00e04c36074c \
-c /etc/suricata/suricata.yaml \
--set outputs.1.eve-log.enabled=yes
```

Suricata processes packets using:

- Signature matching: Rules from Emerging Threats, custom rules
- Protocol analysis: HTTP, DNS, TLS, SSH, etc.
- Flow tracking: Monitors bi-directional connections
- Outputs: EVE JSON format (flows, alerts, HTTP logs, etc.)

#### Step 3: Kafka Bridge

```
# Watches eve.json file and streams to Kafka
with open('/var/log/suricata/eve.json', 'r') as f:
    for line in follow_file(f):
        event = json.loads(line)
        producer.send('suricata-alerts', event)
```

#### The bridge:

- Tails the EVE JSON log file
- Parses each JSON event
- Publishes to Kafka topic
- Handles reconnections and errors

#### Step 4: ML Consumer Processing

```
# Extract features from flow events
features = extract_cicids2017_features(flow_event)
# 65 features: packet counts, byte counts, IAT stats, flags, etc.

# ML inference
prediction = model.predict(features)
# Classes: BENIGN, DDoS, PortScan, BotNet, etc.

# Combine with Suricata alerts
threat_score = calculate_combined_score(ml_prediction, suricata_alerts)
```

#### The ML consumer:

- Extracts 65 CICIDS2017 features from network flows
- Maps to 34 model features (feature engineering)
- Performs ML inference using trained Random Forest/LightGBM
- Combines scores with Suricata signature-based alerts
- Publishes enhanced alerts back to Kafka

#### **III** Feature Extraction

The pipeline extracts **65 network flow features** including:

Category	Features	Examples
Basic	8 features	Source/Dest IP, Port, Protocol, Timestamp
Flow Duration	1 feature	Total flow duration
Packet Stats	14 features	Total fwd/bwd packets, lengths, rates
Byte Stats	8 features	Total bytes, mean packet sizes
Inter-Arrival Time	12 features	IAT mean, std, min, max (fwd/bwd)
Flags	6 features	FIN, SYN, RST, PSH, ACK, URG counts
Header Lengths	4 features	Fwd/Bwd header lengths
Bulk Stats	6 features	Bulk rates, averages
Active/Idle	6 features	Active mean/std, Idle mean/std



The ML models detect:

- DDoS (SYN flood, UDP flood, HTTP flood)
- Port Scanning (Nmap, Masscan patterns)
- **Brute Force** (SSH, FTP, Web login attacks)
- **Botnet** (C&C communication patterns)
- Web Attacks (SQL injection, XSS patterns)
- **Infiltration** (Internal reconnaissance)

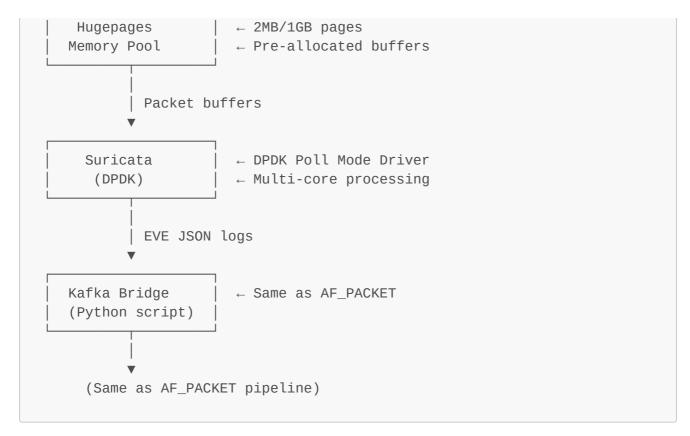
## DPDK Pipeline



Network Traffic

Network Interface ← DPDK-compatible NIC only
(DPDK PMD) ← Intel, Mellanox, Broadcom
← Kernel bypass!

Raw Packets (Zero-copy)



How It Works

#### Step 1: DPDK Setup

#### **Hugepages Allocation:**

```
# Allocate 2GB of 2MB hugepages
echo 1024 > /proc/sys/vm/nr_hugepages
mount -t hugetlbfs nodev /mnt/huge
```

Hugepages provide:

• Reduced TLB misses: Fewer page table lookups

• Contiguous memory: Better for DMA operations

• **Performance**: 10-20% improvement for packet processing

#### **Interface Binding:**

```
# Unbind from kernel driver
dpdk-devbind.py --unbind 0000:01:00.0

# Bind to DPDK driver
dpdk-devbind.py --bind=vfio-pci 0000:01:00.0
```

This removes the NIC from kernel control and gives DPDK direct access.

#### Step 2: DPDK Packet Capture

#### Poll Mode Driver (PMD):

```
Traditional (Interrupt-driven):

Packet arrives → Hardware interrupt → Context switch → Kernel → Copy →
Userspace

DPDK (Polling):

Packet arrives → Direct DMA → Hugepage buffer → Userspace reads
(No interrupts, no context switches, no kernel involvement)
```

#### Benefits:

- Ultra-low latency: < 100 ns packet processing
- **Zero-copy**: Direct memory access to packet buffers
- Multi-core: Dedicated CPU cores for packet processing
- Throughput: 10+ Gbps on commodity hardware

#### Step 3: Suricata DPDK Mode

```
suricata --dpdk \
   -c /etc/suricata/suricata-dpdk.yaml \
   --set dpdk.interfaces.0.interface=0000:01:00.0
```

#### Suricata with DPDK:

- Uses **DPDK libraries** for packet I/O
- Bypasses kernel network stack completely
- Allocates packet buffers from hugepages
- Uses lock-free queues for multi-core processing
- Same detection logic as AF\_PACKET mode (rules, signatures)

#### Step 4-6: Kafka → ML Consumer

Same as AF\_PACKET pipeline - Once packets are processed by Suricata, the rest of the pipeline is identical.

#### ₹ Performance Characteristics

Metric	AF_PACKET	DPDK
Throughput	1-2 Gbps	10+ Gbps
Latency	10-50 μs	< 1 μs
CPU Usage	Moderate	High (dedicated cores)

Metric	AF_PACKET	DPDK
Packet Loss	5-10% @ high load	< 1% @ high load
Setup Complexity	Easy	Complex

## Comparison

When to Use AF\_PACKET Mode

### Use AF\_PACKET if:

- You have a USB network adapter
- You're testing/development
- Traffic < 1 Gbps
- You want simple setup
- You don't have DPDK-compatible hardware
- You need WiFi or virtual interfaces

#### When to Use DPDK Mode

#### **Use DPDK if:**

- You have a DPDK-compatible NIC (Intel X710, Mellanox ConnectX, etc.)
- Traffic > 1 Gbps
- You need minimal packet loss
- Low latency is critical
- You can dedicate CPU cores
- Production deployment

#### Hardware Requirements

Component	AF_PACKET	DPDK
NIC	Any interface	DPDK-compatible only
RAM	4GB+	8GB+ (hugepages)
СРИ	2+ cores	4+ cores (dedicated)
os	Any Linux	Linux with hugepages

## Component Details

#### 1. Suricata

Role: Signature-based intrusion detection

#### **Key Features:**

- 30,000+ community rules (Emerging Threats)
- Protocol parsers: HTTP, DNS, TLS, SSH, SMB, etc.
- Flow tracking: Monitors connection state
- File extraction: Can extract files from traffic
- Performance: Multi-threaded, hardware accelerated

#### Configuration:

```
# suricata.yaml
af-packet:
    - interface: enx00e04c36074c
        threads: 4
        cluster-type: cluster_flow

outputs:
    - eve-log:
        enabled: yes
        filetype: regular
        types:
        - alert
        - http
        - dns
        - flow
```

#### 2. Apache Kafka

**Role**: Message broker and event streaming platform

#### Why Kafka?

- Scalability: Can handle millions of events/sec
- **Durability**: Persists messages to disk
- **Decoupling**: Producers and consumers are independent
- Replay: Can replay historical events
- Multi-consumer: Multiple consumers can read same stream

#### **Topics:**

- suricata-alerts: All Suricata events (flows + alerts)
- ml-predictions: ML-enhanced threat predictions

#### 3. ML Consumer

Role: Machine learning inference and threat scoring

#### Components:

- 1. Feature Extractor (feature\_extractor.py)
  - Parses Suricata flow events

- Calculates 65 CICIDS2017 features
- Handles missing data
- 2. Feature Mapper (feature\_mapper.py)
  - Maps 65 features → 34 model features
  - Handles feature engineering
  - Normalizes values
- Model Loader (model\_loader.py)
  - Loads trained ML models (Random Forest, LightGBM)
  - Manages model versions
  - Performs inference
- 4. Alert Processor (alert\_processor.py)
  - Combines ML predictions with Suricata alerts
  - Calculates composite threat scores
  - Generates enhanced alerts

#### ML Models:

- Random Forest (2017 model): 99.2% accuracy on CICIDS2017
- LightGBM (2018 model): 99.5% accuracy on CICIDS2018
- 4. Kafka Bridge

Role: Stream Suricata logs to Kafka

#### Implementation:

```
def tail_file(filename):
    """Follow a file like 'tail -f'"""
    with open(filename, 'r') as f:
        f.seek(0, 2) # Go to end
        while True:
            line = f.readline()
            if not line:
                time.sleep(0.1)
                continue
            yield line

for line in tail_file('/var/log/suricata/eve.json'):
        event = json.loads(line)
        producer.send('suricata-alerts', event)
```

```
1. Client sends HTTP GET request to suspicious domain
   └ Source: 192.168.1.100:54321
   └ Dest: 192.0.2.50:80
   └─ Payload: "GET /malware.exe HTTP/1.1"
2. Suricata captures packet via AF_PACKET/DPDK
   └─ Signature match: "ET MALWARE Suspicious .exe Download"
   └─ HTTP parser extracts: hostname, URI, user-agent
   └ Flow tracker: This is packet 5 in the flow
3. Suricata writes EVE JSON:
  "timestamp": "2025-10-09T10:15:30.123456+0000",
  "event_type": "alert",
  "alert": {
    "signature": "ET MALWARE Suspicious .exe Download",
    "category": "Malware Command and Control Activity Detected",
    "severity": 1
  },
  "flow": {
    "pkts_toserver": 5,
    "pkts_toclient": 4,
    "bytes_toserver": 480,
    "bytes_toclient": 8420
  },
  "http": {
    "hostname": "bad-domain.com",
    "url": "/malware.exe",
    "http_method": "GET"
 }
}
4. Kafka Bridge reads and publishes to Kafka
5. ML Consumer receives event:
   a. Extracts 65 features from flow data:
      - Total packets: 9
      - Total bytes: 8900
      - Flow duration: 2.5 seconds
      - Average packet size: 988 bytes
      - Forward IAT mean: 0.5s
      - SYN flags: 1, ACK flags: 8
      ... (59 more features)
   b. Maps to 34 model features
   c. ML model predicts:
      - Class: "Web Attack" (0.85 confidence)
      - Anomaly score: 0.92
```

```
d. Combines with Suricata alert:
    - Suricata severity: 1 (High)
    - ML confidence: 0.85
    - Combined threat score: 0.95 (CRITICAL)
e. Publishes enhanced alert:
"timestamp": "2025-10-09T10:15:30.500000+0000",
"source_event": "suricata",
"ml_prediction": {
 "class": "Web Attack",
 "confidence": 0.85,
  "model": "random_forest_2017"
},
"combined_threat_score": 0.95,
"threat_level": "CRITICAL",
"suricata_alert": { ... },
"recommended_action": "BLOCK"
```

### Performance Tuning

#### AF\_PACKET Optimizations

```
# Increase ring buffer size
ethtool -G enx00e04c36074c rx 4096

# Disable offloading for accurate capture
ethtool -K enx00e04c36074c gro off lro off gso off tso off

# Enable promiscuous mode
ip link set enx00e04c36074c promisc on
```

#### **DPDK Optimizations**

```
# CPU isolation (dedicate cores to DPDK)
# Edit /etc/default/grub:
GRUB_CMDLINE_LINUX="isolcpus=2,3,4,5"

# Increase hugepages
echo 2048 > /proc/sys/vm/nr_hugepages

# Use 1GB hugepages for better performance
echo 4 > /sys/kernel/mm/hugepages/hugepages-1048576kB/nr_hugepages
```

### Troubleshooting

#### AF\_PACKET Issues

Problem: No packets captured

```
# Check interface is up and in promiscuous mode
ip link show enx00e04c36074c

# Check Suricata is running
ps aux | grep suricata

# Check Suricata logs
tail -f /var/log/suricata/suricata.log
```

Problem: High packet loss

```
# Increase Suricata threads
suricata --af-packet=enx00e04c36074c --set af-packet.threads=4

# Check CPU usage
top -H -p $(pgrep suricata)
```

#### **DPDK** Issues

Problem: Interface binding fails

```
# Check IOMMU is enabled
dmesg | grep -i iommu

# Try different driver
dpdk-devbind.py --bind=uio_pci_generic 0000:01:00.0
```

**Problem:** Suricata crashes with DPDK

```
# Check hugepages
cat /proc/meminfo | grep Huge

# Check DPDK EAL parameters
suricata --dpdk --dump-config | grep dpdk
```

Summary

Both pipelines provide comprehensive threat detection combining:

- **Signature-based detection** (Suricata rules)
- **Machine Learning** Anomaly detection (Machine Learning)
- V Flow-based analysis (Network behavior)
- **Real-time processing** (Kafka streaming)
- **Scalable architecture** (Microservices)

**Choose AF\_PACKET** for ease of use and compatibility.

**Choose DPDK** for maximum performance and minimal packet loss.

+ 13 / 13 +

Both modes share the same ML detection capabilities and achieve the same detection accuracy!